

```
import pandas as pd
import seaborn as sb
```

```
df = pd.read_csv("/content/student_pass_fail_dataset.csv")
```

```
print(df.head())
```

```

hours_studied  attendance_percentage  assignments_submitted  pass_fail
0             0.5                   35                      0         0
1             1.0                   45                      1         0
2             1.5                   60                      2         0
3             2.0                   70                      3         1
4             2.5                   75                      3         1

```

```
df.describe()
```

```

hours_studied  attendance_percentage  assignments_submitted  pass_fail
count          15.000000           15.000000           15.000000  15.000000
mean           2.400000           67.533333           2.733333   0.666667
std            1.490446           23.311146           1.751190   0.487950
min            0.000000           20.000000           0.000000   0.000000
25%            1.250000           52.500000           1.500000   0.000000
50%            2.000000           70.000000           3.000000   1.000000
75%            3.500000           86.500000           4.000000   1.000000
max            5.000000          100.000000           5.000000   1.000000

```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15 entries, 0 to 14
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   hours_studied          15 non-null    float64
1   attendance_percentage  15 non-null    int64
2   assignments_submitted  15 non-null    int64
3   pass_fail              15 non-null    int64
dtypes: float64(1), int64(3)
memory usage: 612.0 bytes

```

```
df['hours_studied'].unique()
```

```
array([0.5, 1. , 1.5, 2. , 2.5, 3. , 3.5, 4. , 0. , 4.5, 5. ])
```

```
from sklearn.preprocessing import LabelEncoder
```

```
l=LabelEncoder()
```

```
df['hours_studied']=l.fit_transform(df['hours_studied'])
```

```
df['hours_studied'].unique()
```

```
array([ 1,  2,  3,  4,  5,  6,  7,  8,  0,  9, 10])
```

```
df
```



	hours_studied	attendance_percentage	assignments_submitted	pass_fail
0	1	35	0	0
1	2	45	1	0
2	3	60	2	0
3	4	70	3	1
4	5	75	3	1
5	6	80	4	1
6	7	85	4	1
7	8	90	5	1
8	0	20	0	0
9	9	95	5	1
10	10	100	5	1
11	2	50	1	0
12	4	65	2	1
13	4	55	2	1
14	7	88	4	1

```
from sklearn.model_selection import train_test_split
```

```
X = df[['hours_studied', 'attendance_percentage', 'assignments_submitted']]
```

```
y = df['pass_fail']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, )
```


```
sb.pairplot(df,hue='hours_studied')
```

```
<seaborn.axisgrid.PairGrid at 0x7d1653867a90>
```



```
from sklearn.linear_model import LogisticRegression
```

```
lr = LogisticRegression()
lr.fit(X_train, y_train)
```



```
y_pred = lr.predict(X_test)
```

```
from sklearn.metrics import accuracy_score
```

```
accuracy_score(y_test, y_pred)
```



```
import pickle
```

```
pickle.dump(lr, open('students.pkl', 'wb'))
```

Start coding or [generate](#) with AI.

